



ÉCOLE POLYTECHNIQUE  
FÉDÉRALE DE LAUSANNE

# Reduction Operations

Big Data Analysis with Scala and Spark

Heather Miller

# What we've seen so far

- ▶ we defined *Distributed Data Parallelism*
- ▶ we saw that Apache Spark implements this model
- ▶ we got a feel for what latency means to distributed systems

# What we've seen so far

- ▶ we defined *Distributed Data Parallelism*
- ▶ we saw that Apache Spark implements this model
- ▶ we got a feel for what latency means to distributed systems

## Spark's Programming Model

- ▶ We saw that, at a glance, Spark looks like Scala collections
- ▶ However, internally, Spark behaves differently than Scala collections
  - ▶ Spark uses ***laziness*** to save time and memory
- ▶ We saw *transformations* and *actions*
- ▶ We saw caching and persistence (*i.e.*, cache in memory, save time!)
- ▶ We saw how the cluster topology comes into the programming model

# Transformations to Actions

Most of our intuitions have focused on distributing **transformations** such as `map`, `flatMap`, `filter`, etc.

*We've visualized how transformations like these are distributed and parallelized.*

# Transformations to Actions

Most of our intuitions have focused on distributing **transformations** such as map, flatMap, filter, etc.

*We've visualized how transformations like these are distributed and parallelized.*

**But what about actions? In particular, how are common reduce-like actions distributed in Spark?**

# Reduction Operations, Generally

**First, what do we mean by “reduction operations”?**

Recall operations such as `fold`, `reduce`, and `aggregate` from Scala sequential collections. All of these operations and their variants (such as `foldLeft`, `reduceRight`, etc) have something in common.

# Reduction Operations, Generally

**First, what do we mean by “reduction operations”?**

Recall operations such as fold, reduce, and aggregate from Scala sequential collections. All of these operations and their variants (such as foldLeft, reduceRight, etc) have something in common.

**Reduction Operations:**

**walk through a collection and combine neighboring elements of the collection together to produce a single combined result.**

*(rather than another collection)*

# Reduction Operations, Generally

## Reduction Operations:

walk through a collection and combine neighboring elements of the collection together to produce a single combined result.

*(rather than another collection)*

## Example:

```
case class Taco(kind: String, price: Double)
```

```
val tacoOrder =  
  List(  
    Taco("Carnitas", 2.25),  
    Taco("Corn", 1.75),  
    Taco("Barbacoa", 2.50),  
    Taco("Chicken", 2.00))
```

```
val cost = tacoOrder.foldLeft(0.0)((sum, taco) => sum + taco.price)
```



# Parallel Reduction Operations

**Recall what we learned in the course Parallel Programming course about foldLeft vs fold.**

Which of these two were parallelizable?

# Parallel Reduction Operations

Recall what we learned in the course **Parallel Programming** course about `foldLeft` vs `fold`.


Which of these two were parallelizable?

**`foldLeft` is not parallelizable.**

```
def foldLeft[B](z: B)(f: (B, A) => B): B
```

*Applies a binary operator to a start value and all elements of this collection or iterator, going left to right.*

— Scala API documentation



# Parallel Reduction Operations: FoldLeft

**foldLeft is not parallelizable.**

```
def foldLeft[B](z: B)(f: (B, A) => B): B
```

Being able to change the result type from A to B forces us to have to execute foldLeft sequentially from left to right.

Concretely, given:

*"1234"*

```
val xs = List(1, 2, 3, 4)
```

```
val res = xs.foldLeft("")(str: String, i: Int) => str + i)
```

What happens if we try to break this collection in two and parallelize?

# Parallel Reduction Operations: FoldLeft

**foldLeft is not parallelizable.**

```
def foldLeft[B](z: B)(f: (B, A) => B): B
```

```
val xs = List(1, 2, 3, 4)
```

```
val res = xs.foldLeft("")(str: String, i: Int) => str + i
```

List(1, 2)

" " + 1 → "1"  
"1" + 2 → "12"  
string

List(3, 4)

" " + 3 → "3"  
"3" + 4 → "34"  
String

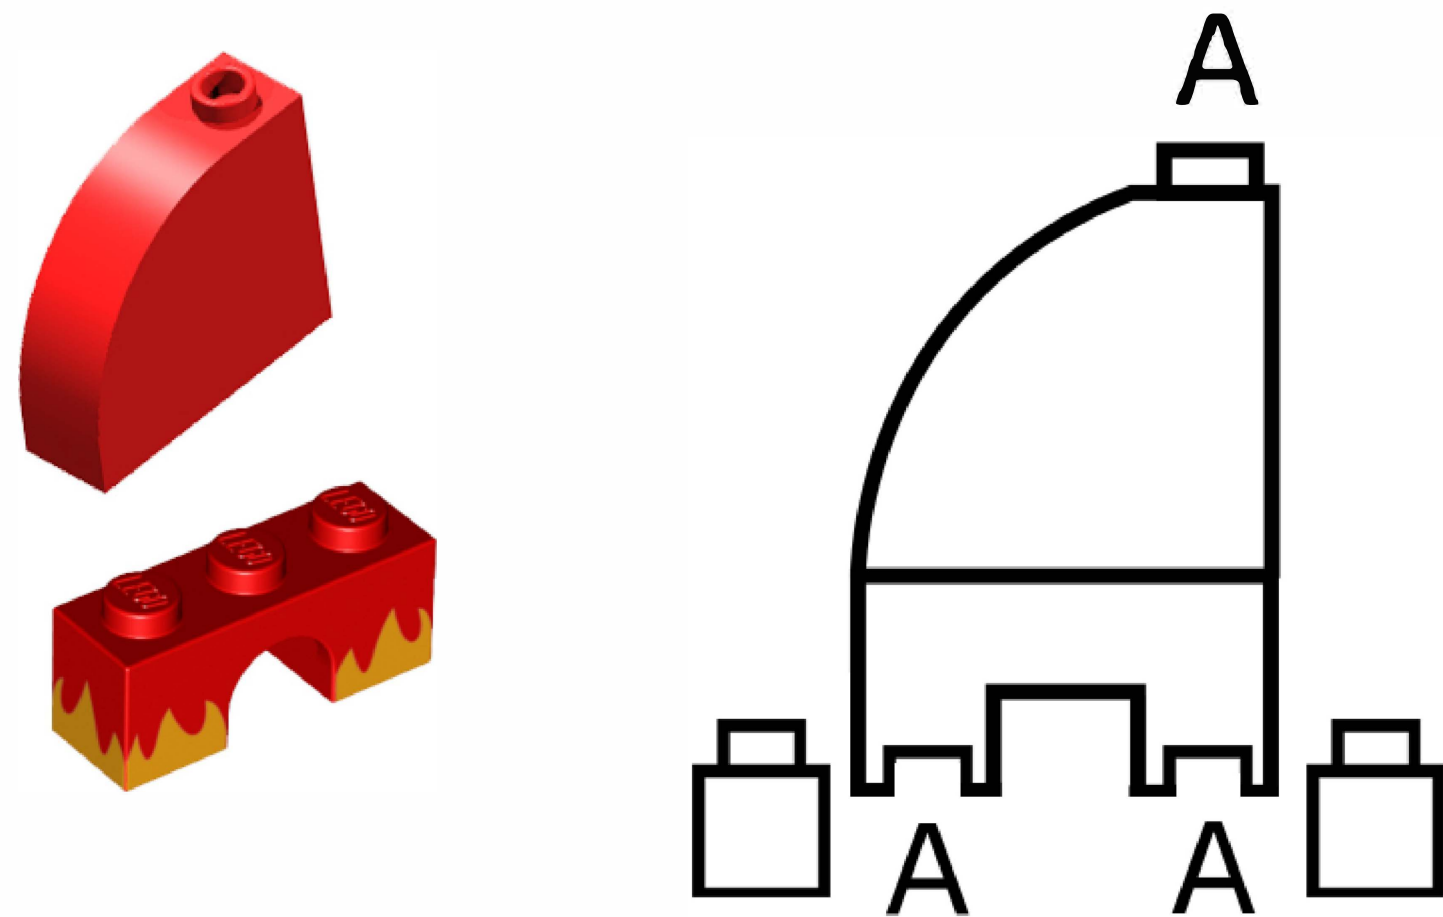
!!! type error !!!

can't apply  
(str: String, i: Int) => str + i !!!

# Parallel Reduction Operations: Fold

fold enables us to parallelize things, but it restricts us to always returning the same type.

```
def fold(z: A)(f: (A, A) => A): A
```

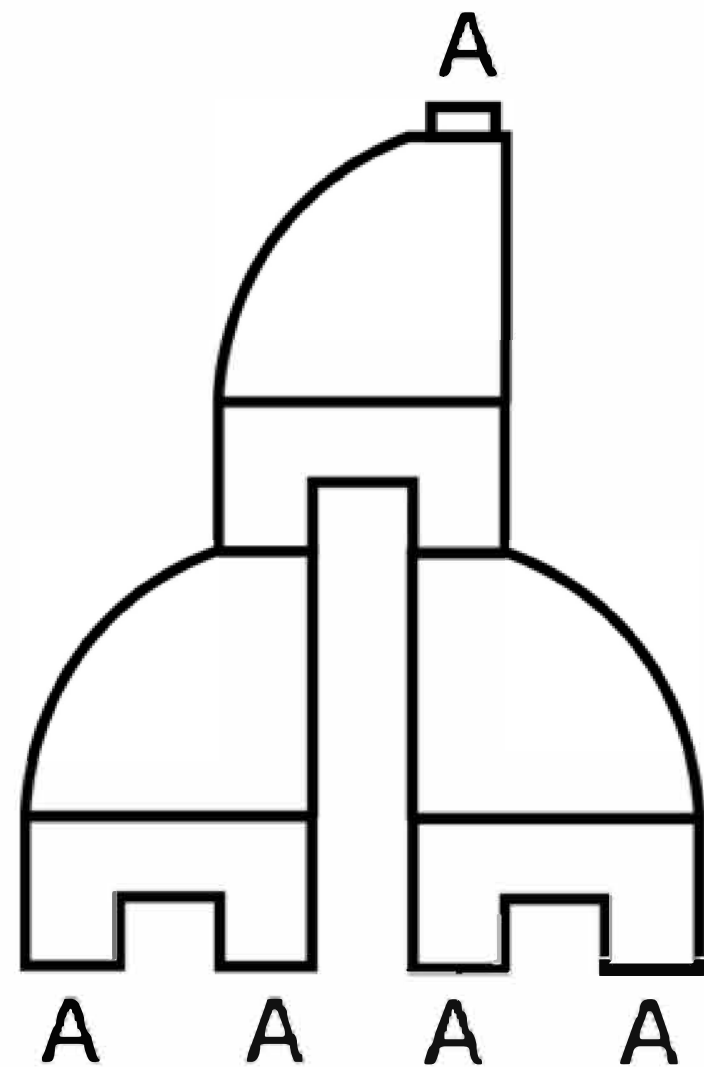



It enables us to parallelize using a single function  $f$  by enabling us to build parallelizable reduce trees.

# Parallel Reduction Operations: Fold

It enables us to parallelize using a single function  $f$  by enabling us to build parallelizable reduce trees.

```
def fold(z: A)(f: (A, A) => A): A
```



# Parallel Reduction Operations: Aggregate

Does anyone remember what aggregate does?

# Parallel Reduction Operations: Aggregate

Does anyone remember what aggregate does?

```
aggregate[B](z: => B)(seqop: (B, A) => B, combop: (B, B) => B): B
```



# Parallel Reduction Operations: Aggregate

Does anyone remember what aggregate does?

```
aggregate[B](z: => B)(seqop: (B, A) => B, combop: (B, B) => B): B
```


aggregate is said to be general because it gets you the best of both worlds.

## Properties of aggregate

1. Parallelizable.
2. Possible to change the return type.

# Parallel Reduction Operations: Aggregate

`aggregate[B](z: => B)(seqop: (B, A) => B, combop: (B, B) => B): B`



Aggregate lets you still do sequential-style folds *in chunks* which change the result type. Additionally requiring the combop function enables building one of these nice reduce trees that we saw is possible with fold to *combine these chunks* in parallel.

# Reduction Operations on RDDs

## Scala collections:

fold

foldLeft/foldRight

reduce

aggregate

## Spark:

fold

~~foldLeft/foldRight~~

reduce

aggregate

# Reduction Operations on RDDs

## Scala collections:

fold

foldLeft/foldRight

reduce

aggregate

## Spark:

fold

~~foldLeft/foldRight~~

reduce

aggregate

Spark doesn't even give you the option to use foldLeft/foldRight. Which means that if you have to change the return type of your reduction operation, your only choice is to use aggregate.

# Reduction Operations on RDDs

## Scala collections:

fold

foldLeft/foldRight

reduce

aggregate

## Spark:

fold

~~foldLeft/foldRight~~

reduce

aggregate

Spark doesn't even give you the option to use foldLeft/foldRight. Which means that if you have to change the return type of your reduction operation, your only choice is to use aggregate.

***Question:** Why not still have a serial foldLeft/foldRight on Spark?*

# Reduction Operations on RDDs

## Scala collections:

fold

foldLeft/foldRight

reduce

aggregate

## Spark:

fold

~~foldLeft/foldRight~~

reduce

aggregate

Spark doesn't even give you the option to use foldLeft/foldRight. Which means that if you have to change the return type of your reduction operation, your only choice is to use aggregate.

***Question:** Why not still have a serial foldLeft/foldRight on Spark?*

*Doing things serially across a cluster is actually difficult. Lots of synchronization. Doesn't make a lot of sense.*

# RDD Reduction Operations: Aggregate

In Spark, aggregate is a more desirable reduction operator a majority of the time. Why do you think that's the case?

# RDD Reduction Operations: Aggregate

In Spark, aggregate is a more desirable reduction operator a majority of the time. Why do you think that's the case?

As you will realize from experimenting with our Spark <sup>assignments</sup>~~cluster~~, much of the time when working with large-scale data, our goal is to ***project down from larger/more complex data types.***




# RDD Reduction Operations: Aggregate

In Spark, aggregate is a more desirable reduction operator a majority of the time. Why do you think that's the case?

As you will realize from experimenting with our Spark cluster, much of the time when working with large-scale data, our goal is to ***project down from larger/more complex data types.***

**Example:**

```
case class WikipediaPage(  
  title: String,  
  redirectTitle: String,  
  timestamp: String,  
  lastContributorUsername: String,  
  text: String)
```



## RDD Reduction Operations: Aggregate

As you will realize after experimenting with Spark a bit, much of the time when working with large-scale data, your goal is to ***project down from larger/more complex data types***.

**Example:**

```
case class WikipediaPage(  
  title: String,  
  redirectTitle: String,  
  timestamp: String,  
  lastContributorUsername: String,  
  text: String)
```

I might only care about title and timestamp, for example. In this case, it'd save a lot of time/memory to not have to carry around the full-text of each article (text) in our accumulator!

**Hence, why accumulate is often more desirable in Spark than in Scala collections!**